

# Using the Intel Distribution of the OpenVINO Toolkit for Deploying Accelerated Deep Learning Applications [2021.1]



# Notices and disclaimer

- INFORMATION IN THIS DOCUMENT IS PROVIDED IN CONNECTION WITH INTEL PRODUCTS. NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. EXCEPT AS PROVIDED IN INTEL'S TERMS AND CONDITIONS OF SALE FOR SUCH PRODUCTS, INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO SALE AND/OR USE OF INTEL PRODUCTS INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT OR OTHER INTELLECTUAL PROPERTY RIGHT.
- Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Performance varies depending on system configuration. No computer system can be absolutely secure. Check with your system manufacturer or retailer or learn more at [www.intel.com](http://www.intel.com).
- This document contains information on products, services and/or processes in development. All information provided here is subject to change without notice. Contact your Intel representative to obtain the latest forecast, schedule, specifications and roadmaps.
- © Intel Corporation. Intel, the Intel logo, and other Intel marks are trademarks of Intel Corporation or its subsidiaries. Other names and brands may be claimed as the property of others.

# Optimization notice

Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Performance varies depending on system configuration. No computer system can be absolutely secure. Check with your system manufacturer to learn more.

The benchmark results reported in this deck may need to be revised as additional testing is conducted. The results depend on the specific platform configurations and workloads utilized in the testing and may not be applicable to any particular user's components, computer system or workloads. The results are not necessarily representative of other benchmarks and other benchmark results may show greater or lesser impact from mitigations.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products. For more complete information visit [www.intel.com/benchmarks](http://www.intel.com/benchmarks).

Cost reduction scenarios described are intended as examples of how a given Intel- based product, in the specified circumstances and configurations, may affect future costs and provide cost savings. Circumstances will vary. Intel does not guarantee any costs or cost reduction.

Other names and brands may be claimed as the property of others. Any third-party information referenced on this document is provided for information only. Intel does not endorse any specific third-party product or entity mentioned on this document. Intel, the Intel logo, and other Intel marks are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries. Copyright Intel Corporation.

## Optimization Notice



Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Notice revision #20110804

OpenVX and the OpenVX logo are trademarks of the Khronos Group Inc.

OpenCL and the OpenCL logo are trademarks of Apple Inc. used by permission by Khronos

# AI CHANGING AND ENABLING EVERY INDUSTRY



AI software market is projected to reach USD 126.0 billion in annual worldwide revenue by 2025<sup>1</sup>



Deep learning software revenue is estimated to grow to USD 67.2 billion by 2025<sup>2</sup>



Global deep learning chip market is expected to reach USD 29.4 billion by 2025<sup>3</sup>

## AGRICULTURE

Achieve higher yields and increase efficiency

## ENERGY

Maximize production and uptime

## EDUCATION

Transform the learning experience

## GOVERNMENT

Enhance safety, research, and more

## FINANCE

Turn data into valuable intelligence

## HEALTH

Revolutionize patient outcomes

## INDUSTRIAL

Empower truly intelligent Industry 4.0

## MEDIA

Create thrilling experiences

## RETAIL

Transform stores and inventory

## SMART HOME

Enable homes that see, hear, and respond

## TELECOM

Drive network and operational efficiency

## TRANSPORTATION

Automated driving

1. Tractica, [Artificial Intelligence Software Market](#), 2020


2. Tractica, [deep learning research](#), 2018

3. AlliedMarketResearch, [Deep Learning Chip Market](#), 2018


# Deep learning breakthroughs and opportunities

Machines able to meet or exceed human image and speech recognition

## Image Recognition



## Speech Recognition



Source: ILSVRC ImageNet winning entry classification error rate each year 2010-2016 (Left), <https://www.microsoft.com/en-us/research/blog/microsoft-researchers-achieve-new-conversational-speech-recognition-milestone/> (Right)

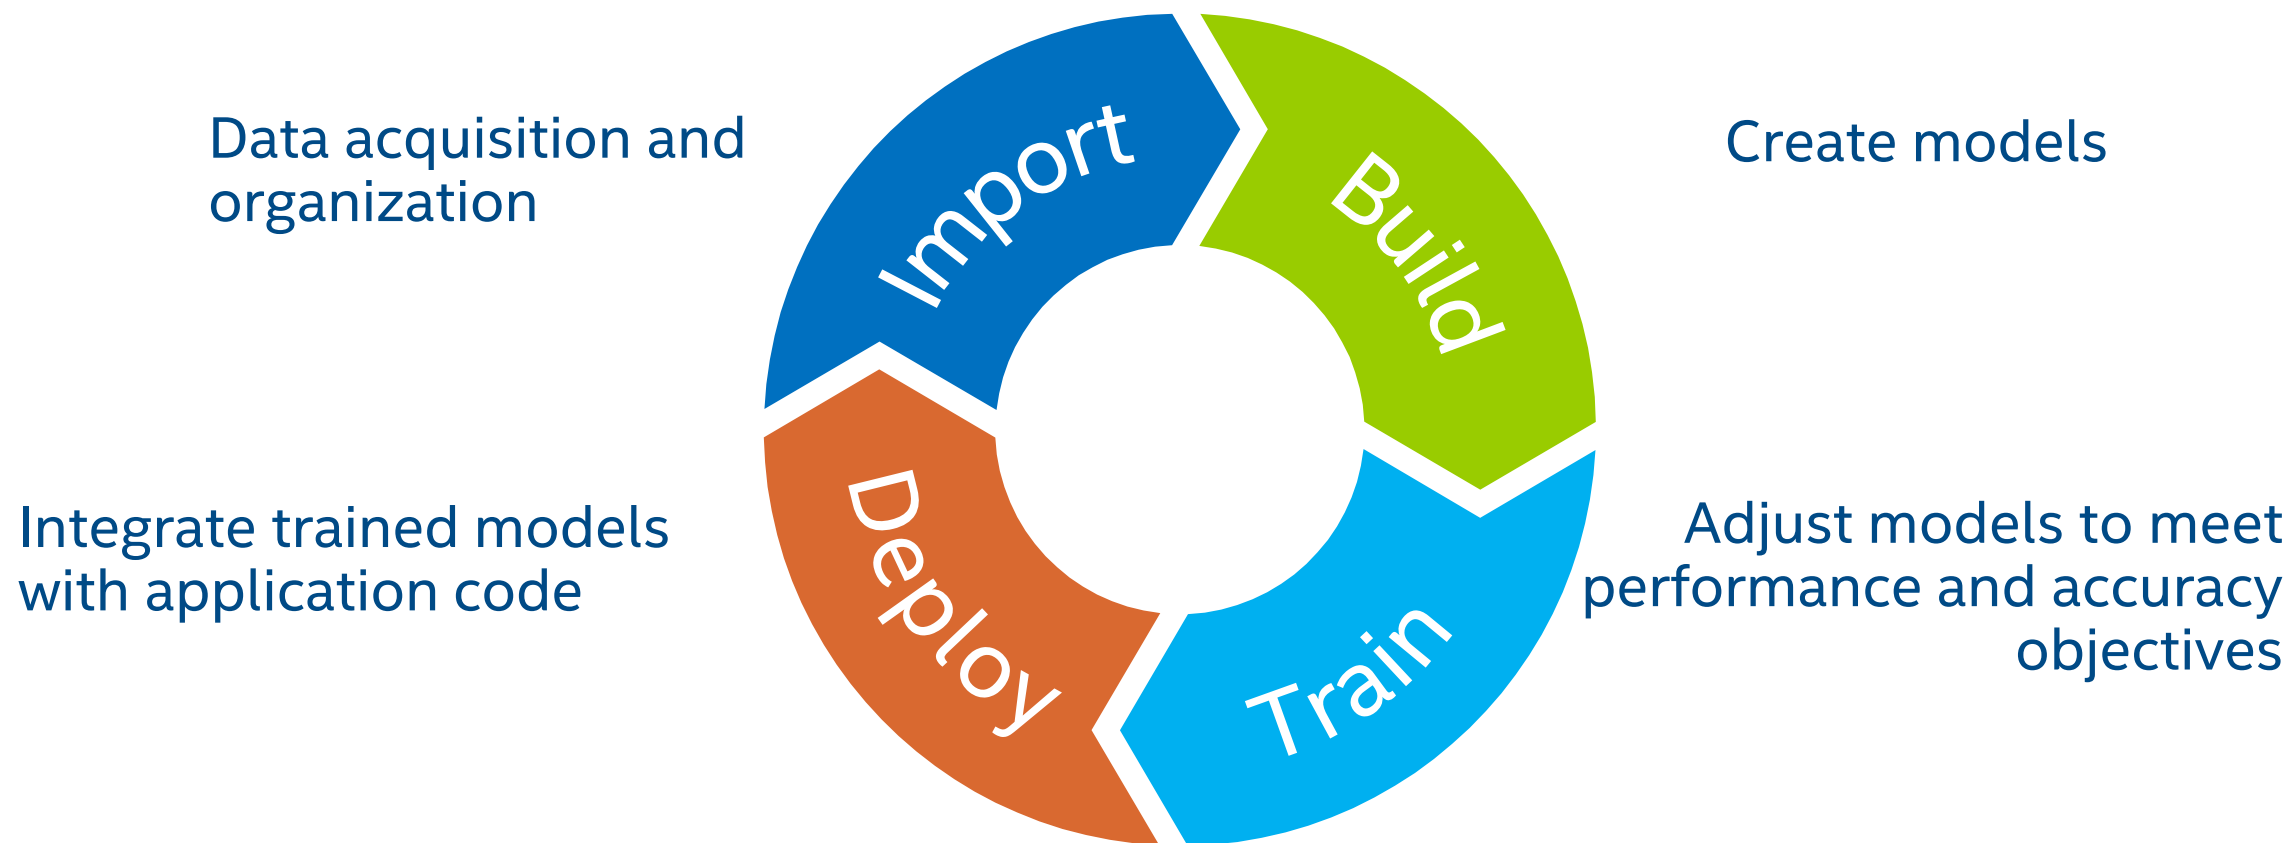
Source: <https://www.mckinsey.com/featured-insights/artificial-intelligence/notes-from-the-ai-frontier-applications-and-value-of-deep-learning>



**ADDITIONAL ECONOMIC  
IMPACT DRIVEN BY AI**

**\$13 TRILLION IN 2030**


# Deep Learning Development Cycle




Intel® Distribution OpenVINO™ Toolkit Provides Deployment from Intel® Edge to Cloud

# Deep Learning: Training vs. Inference

## Training




## Inference



## Did You Know?


Training requires a very large data set and deep neural network (many layers) to achieve the highest accuracy in most cases



# AI COMPUTE CONSIDERATIONS

How do you determine the right computing for your AI needs?

## WORKLOADS




What is my workload profile?

## REQUIREMENTS



What are my use case requirements?

## DEMAND




How prevalent is AI in my environment?



# WHY INTEL AI COMPUTE?

## MAXIMIZE



Get the most out of the foundation for AI from the CPU leader

## OPTIMIZE



Choose the right compute for you from the one with all the options

## SIMPLIFY

OPTIMIZED SW

DATA PIPELINE

ANALYTICS & AI

SUPPORT

MOVE/STORE

Reduce “moving parts” by building on an optimized AI platform



## LEAD




Lead your industry by aligning with the builder of next-gen AI solutions

# Intel® distribution of OpenVINO™ toolkit

- Tool Suite for High-Performance, Deep Learning Inference
- Fast, accurate real-world results using high-performance, AI and computer vision inference deployed into production across Intel® architecture from edge to cloud



High-Performance,  
Deep Learning Inference



Streamlined Development,  
Ease of Use

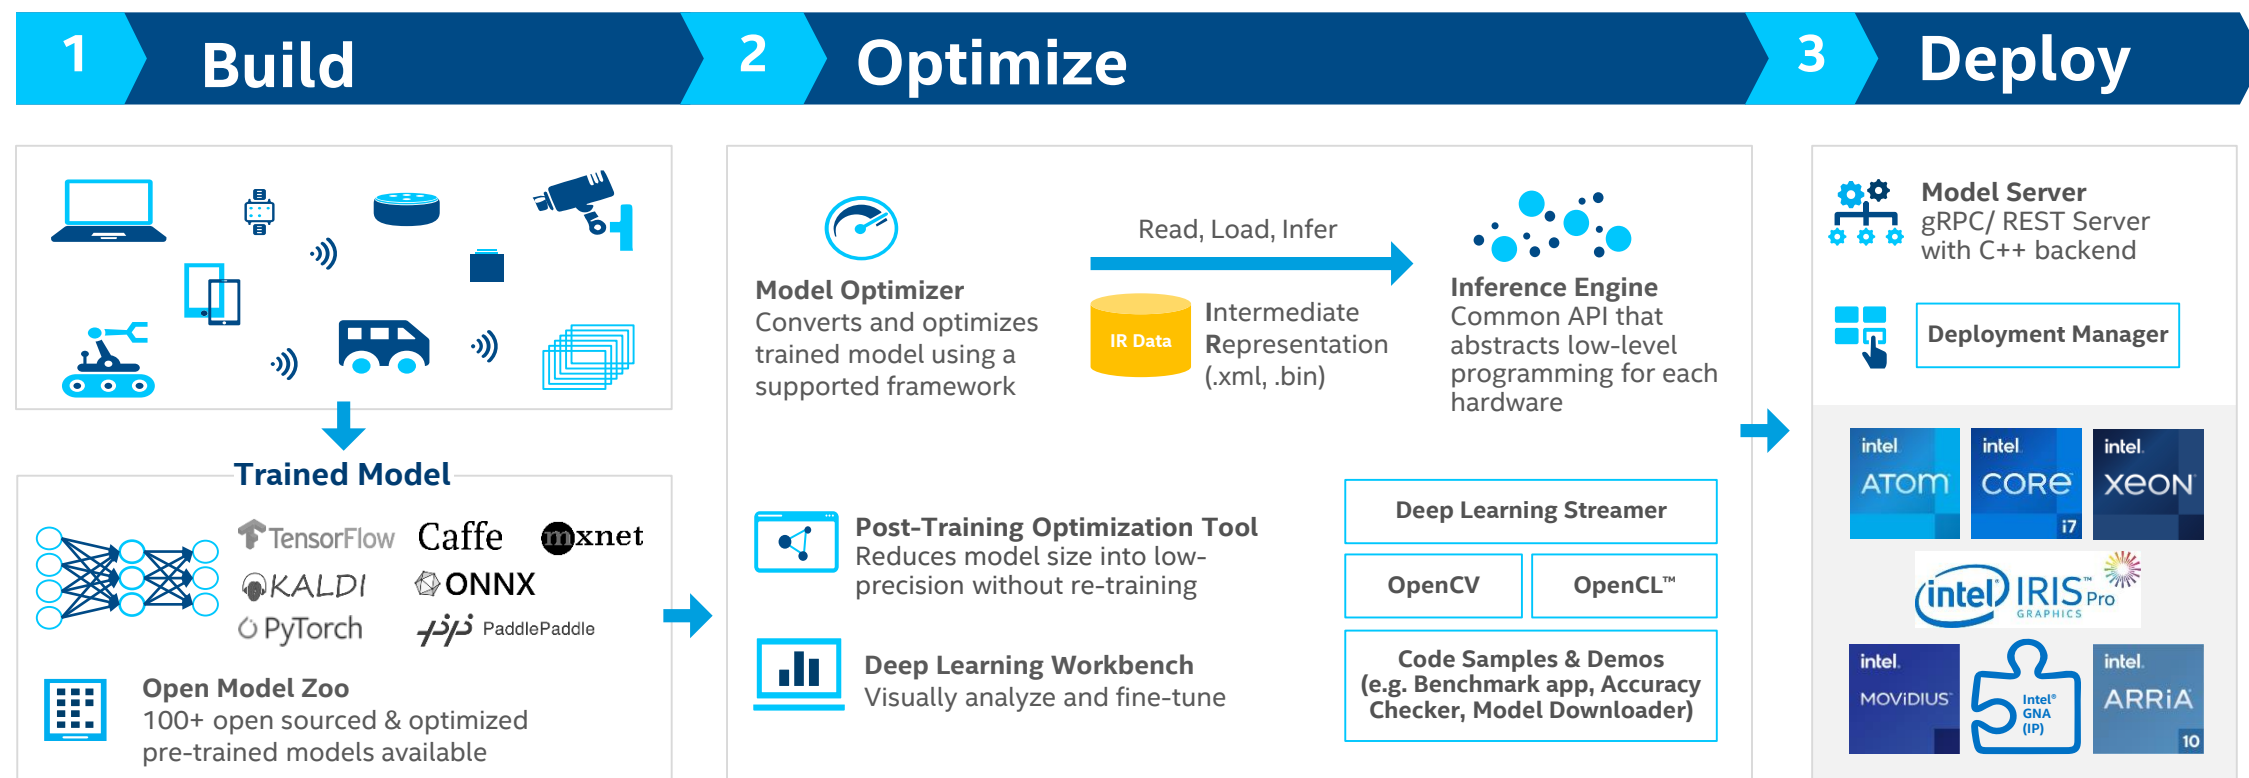


Write Once,  
Deploy Anywhere

## New Features from OpenVINO Toolkit 2021.1

- Support for Tiger Lake (11th generation Intel® Core™ processors)
- New capabilities in OpenVINO™ Model Server
- Support for TensorFlow 2.x
- Support for non-computer vision workloads
- (Coming in Q4) Beta Release: Integration of OpenVINO™ toolkit DL Workbench and Intel® DevCloud for the Edge
- Support for GNA 2.0

# Three steps for the Intel® Distribution of OpenVINO™ toolkit



# Additional Tools and Add-ons from the OpenVINO GitHub Repo

## [Computer Vision Annotation Tool](#)

This web-based tool helps annotate videos and images before training a model

## [Dataset Management Framework](#)

Use this add-on to build, transform and analyze datasets

## [Neural Network Compression Framework](#)

Training framework based on PyTorch\* for quantization-aware training

## *[NEW]* [OpenVINO™ Model Server](#)

Scalable inference server for serving optimized models and applications

## [Training Extensions](#)

Trainable deep learning models for training with custom data

# Speed up development with open source resources

## Open source resources with pre-trained models, samples and demos



### Computer Vision

[Object detection](#)  
[Object recognition](#)  
[Reidentification](#)  
Volumetric segmentation  
[Semantic segmentation](#)  
[Instance segmentation](#)  
3D reconstruction  
[Human pose estimation](#)  
[Image processing](#)  
[Action recognition](#)  
Image super resolution



### Audio, Speech, Language

Language processing  
Speech to text  
[Text detection](#)  
[Text recognition](#)  
Natural Language Processing



### Other

(Data Generation,  
Reinforcement Learning)

[Compressed models](#)  
[Image retrieval](#)



Model  
Downloader



Accuracy  
Checker

- Provides an easy way of accessing a number of public models as well as a set of pre-trained Intel models
- Check for accuracy of the model (original and after conversion) to IR file using a known data set

And more..

## PRE-TRAINED MODELS

[https://github.com/opencv/open\\_model\\_zoo](https://github.com/opencv/open_model_zoo)

# Speed up development with open source resources

## Open source resources with pre-trained models, demos, and tools

The Open Model Zoo demo applications are console applications that demonstrate how you can use your applications to solve specific use-cases.




### Smart Classroom

Recognition and action detection demo for classroom settings



### Weld Porosity Detection

Demonstrates how to find defects in welding



### Multi-Camera, Multi-Person

Tracking multiple people on multiple cameras for public safety use cases



### Person Inpainting

Removes unwanted people in images or videos



### Gaze Estimation

Face detection followed by gaze estimation, head pose estimation and facial landmarks regression.

*And more..*

## DEMO APPLICATIONS

[https://github.com/opencv/open\\_model\\_zoo](https://github.com/opencv/open_model_zoo)

# Choose between Release Types

## Standard Releases vs Long-Term Support Releases



**Standard Release (3-4 releases a year):** Users looking to take advantage of new features, tools and support in order to keep current with the advancements in deep learning technologies



**Long-Term Support Release:** Users looking for a stable and reliable version that is maintained for a longer period of time, and are looking for little to no new feature changes

# Supported OS and Install Options [2021.1]

<https://software.intel.com/content/www/us/en/develop/tools/opencv-toolkit.html>

## ■ Operating Systems

- Ubuntu 18.04.x long-term support (LTS), 64-bit
- CentOS 7.6, 64-bit (for target only)
- Yocto Project v3.0, 64-bit (for target only and requires modifications)
- Microsoft Windows\* 10 64-bit
- macOS\* 10.15
- Raspbian\* Buster, Stretch

## ■ Install From Images and Repositories

- GitHub
  - <https://github.com/opencv-toolkit/opencv.git>
- Anaconda Cloud
  - <https://anaconda.org/intel/opencv-ie4py>

## • Python\* Package Installer (PIP)

- <https://pypi.org/project/opencv-python/>

## • Docker

- [Install from Image file](#)
- [Download from DockerHub](#)

## • APT

- `sudo apt-cache search install-opencv-runtime-ubuntu18`

## • YUM

- `sudo install intel-opencv-runtime-centos7`

## ■ Intel® Edge Software Hub

- [Edge Insights for Vision](#)